

# Association Rules and Clustering

Steve Avsec

Illinois Institute of Technology

April 15, 2024

# Overview

- 1 One Last DL Note
- 2 Association Rules
- 3 Min Hashing

# One Last Note on Convexity

Just as a point of clarification:

- Non-convexity of the loss function in DL is not a “bad” thing.

# One Last Note on Convexity

Just as a point of clarification:

- Non-convexity of the loss function in DL is not a “bad” thing.
- Convex functions are a small, special class of functions that are easy to compute with.

# One Last Note on Convexity

Just as a point of clarification:

- Non-convexity of the loss function in DL is not a “bad” thing.
- Convex functions are a small, special class of functions that are easy to compute with.
- Many interesting physical systems are studied using non-convex optimization problems.

# One Last Note on Convexity

Just as a point of clarification:

- Non-convexity of the loss function in DL is not a “bad” thing.
- Convex functions are a small, special class of functions that are easy to compute with.
- Many interesting physical systems are studied using non-convex optimization problems.
- The class of non-convex optimization functions is NP-complete.

# The Setup

- We have a set of "items"  $\{i_1, \dots, i_L\}$ .

# The Setup

- We have a set of "items"  $\{i_1, \dots, i_L\}$ .
- We have database of "transactions"  $\{X_1, \dots, X_N\}$ .



# The Setup

- We have a set of "items"  $\{i_1, \dots, i_L\}$ .
- We have database of "transactions"  $\{X_1, \dots, X_N\}$ .
- A "typical" setup is that each  $X_j$  consists of a unique id plus a  $k$ -dimensional vector describing each item in the transaction.

# The Setup

- We have a set of "items"  $\{i_1, \dots, i_L\}$ .
- We have database of "transactions"  $\{X_1, \dots, X_N\}$ .
- A "typical" setup is that each  $X_j$  consists of a unique id plus a  $k$ -dimensional vector describing each item in the transaction.
- The basic goal is to estimate  $P(v)$  where  $v$  is a particular feature vector.

## Item Sets

- Let  $S_i$  denote the set of values each feature vector can take.

## Item Sets

- Let  $S_l$  denote the set of values each feature vector can take.
- Let

$$K = \sum_{l=1}^k |S_l|$$

## Item Sets

- Let  $S_l$  denote the set of values each feature vector can take.
- Let

$$K = \sum_{l=1}^k |S_l|$$

- Let  $Z_k$  denote a dummy variable  $\{0, 1\}$ -valued that indicates if the corresponding item takes a particular value.

## Item Sets

- Let  $S_l$  denote the set of values each feature vector can take.
- Let

$$K = \sum_{l=1}^k |S_l|$$

- Let  $Z_k$  denote a dummy variable  $\{0, 1\}$ -valued that indicates if the corresponding item takes a particular value.
- Goal: Find subsets  $\mathcal{K} \subset \{1, \dots, K\}$  such that

$$Pr(\cap_{k \in \mathcal{K}} Z_k = 1) = Pr\left(\prod_{k \in \mathcal{K}} Z_k = 1\right)$$

is "large".

# Prevalence

Consider the empirical probability

$$\hat{Pr} \left( \prod_{k \in \mathcal{K}} Z_k = 1 \right) = \frac{1}{N} \sum_{j=1}^N \prod_{k \in \mathcal{K}} z_{j,k}$$

where  $z_{j,k}$  is the value of  $Z_k$  for the  $j$ th transaction.

# Prevalence

Consider the empirical probability

$$\hat{Pr} \left( \prod_{k \in \mathcal{K}} z_k = 1 \right) = \frac{1}{N} \sum_{j=1}^N \prod_{k \in \mathcal{K}} z_{j,k}$$

where  $z_{j,k}$  is the value of  $Z_k$  for the  $j$ th transaction.

This quantity is called the "prevalence" of  $\mathcal{K}$  and is denoted  $\mathcal{T}(\mathcal{K})$ .



# Prevalence

Consider the empirical probability

$$\hat{Pr} \left( \prod_{k \in \mathcal{K}} z_k = 1 \right) = \frac{1}{N} \sum_{j=1}^N \prod_{k \in \mathcal{K}} z_{j,k}$$

where  $z_{j,k}$  is the value of  $Z_k$  for the  $j$ th transaction.

This quantity is called the "prevalence" of  $\mathcal{K}$  and is denoted  $\mathcal{T}(\mathcal{K})$ .

We will look for all item sets such that

$$\{\mathcal{K} | \mathcal{T}(\mathcal{K}) > t\}$$

for some fixed threshold  $t$ .

# The A Priori Algorithm

In "typical", real world data,

- 1 The set

$$\{\mathcal{K} | \mathcal{T}(\mathcal{K}) > t\}$$

is relatively small.

- 2 If  $\mathcal{L} \subseteq \mathcal{K}$  then  $\mathcal{T}(\mathcal{L}) \leq \mathcal{T}(\mathcal{K})$ .

# The A Priori Algorithm

In “typical”, real world data,

- 1 The set

$$\{\mathcal{K} | \mathcal{T}(\mathcal{K}) > t\}$$

is relatively small.

- 2 If  $\mathcal{L} \subseteq \mathcal{K}$  then  $\mathcal{T}(\mathcal{L}) \leq \mathcal{T}(\mathcal{K})$ .

First, compute the prevalence for all single-item sets and keep any items where the prevalence is greater than  $t$ .

# The A Priori Algorithm

In “typical”, real world data,

- 1 The set

$$\{\mathcal{K} | \mathcal{T}(\mathcal{K}) > t\}$$

is relatively small.

- 2 If  $\mathcal{L} \subseteq \mathcal{K}$  then  $\mathcal{T}(\mathcal{L}) \leq \mathcal{T}(\mathcal{K})$ .

First, compute the prevalence for all single-item sets and keep any items where the prevalence is greater than  $t$ .

Then compute the prevalence for all two-item sets where each item passed the previous iteration.

# The A Priori Algorithm

In “typical”, real world data,

- 1 The set

$$\{\mathcal{K} | \mathcal{T}(\mathcal{K}) > t\}$$

is relatively small.

- 2 If  $\mathcal{L} \subseteq \mathcal{K}$  then  $\mathcal{T}(\mathcal{L}) \leq \mathcal{T}(\mathcal{K})$ .

First, compute the prevalence for all single-item sets and keep any items where the prevalence is greater than  $t$ .

Then compute the prevalence for all two-item sets where each item passed the previous iteration.

Continue until all sets have prevalence less than the threshold.

# Association Rules

For a specific subset  $\mathcal{K}$  such that  $\mathcal{T}(\mathcal{K}) > t$ , an *association rule* is just a splitting

$$A \cup B = \mathcal{K} \text{ and } A \cap B = \emptyset$$

# Association Rules

For a specific subset  $\mathcal{K}$  such that  $\mathcal{T}(\mathcal{K}) > t$ , an *association rule* is just a splitting

$$A \cup B = \mathcal{K} \text{ and } A \cap B = \emptyset$$

In this case, we write  $A \Rightarrow B$ , and denote  $T(A \Rightarrow B) = \mathcal{T}(\mathcal{K})$ .

# Association Rules

For a specific subset  $\mathcal{K}$  such that  $\mathcal{T}(\mathcal{K}) > t$ , an *association rule* is just a splitting

$$A \cup B = \mathcal{K} \text{ and } A \cap B = \emptyset$$

In this case, we write  $A \Rightarrow B$ , and denote  $T(A \Rightarrow B) = \mathcal{T}(\mathcal{K})$ . We define the *confidence*  $C(A \Rightarrow B)$  of the rule by

$$C(A \Rightarrow B) = \frac{T(A \Rightarrow B)}{T(A)}$$

which is just an estimate of  $Pr(B|A)$ .



# Association Rules

For a specific subset  $\mathcal{K}$  such that  $\mathcal{T}(\mathcal{K}) > t$ , an *association rule* is just a splitting

$$A \cup B = \mathcal{K} \text{ and } A \cap B = \emptyset$$

In this case, we write  $A \Rightarrow B$ , and denote  $T(A \Rightarrow B) = \mathcal{T}(\mathcal{K})$ . We define the *confidence*  $C(A \Rightarrow B)$  of the rule by

$$C(A \Rightarrow B) = \frac{T(A \Rightarrow B)}{T(A)}$$

which is just an estimate of  $Pr(B|A)$ . Define the *lift*  $L(A \Rightarrow B)$  by

$$L(A \Rightarrow B) = \frac{C(A \Rightarrow B)}{T(B)}$$

# Output

We want to find association rules  $A \Rightarrow B$  such that

- 1  $T(A \Rightarrow B) > t$
- 2  $C(A \Rightarrow B) > c$

# Output

We want to find association rules  $A \Rightarrow B$  such that

- 1  $T(A \Rightarrow B) > t$
- 2  $C(A \Rightarrow B) > c$

The original A Priori algorithm paper (Agrawal et. al. 1995) provides a technique for “pruning” rules for a given subset  $\mathcal{K}$  so one doesn’t have to just check all of the subsets to optimize  $C(A \Rightarrow B)$ .

# Output

We want to find association rules  $A \Rightarrow B$  such that

- 1  $T(A \Rightarrow B) > t$
- 2  $C(A \Rightarrow B) > c$

The original A Priori algorithm paper (Agrawal et. al. 1995) provides a technique for "pruning" rules for a given subset  $\mathcal{K}$  so one doesn't have to just check all of the subsets to optimize  $C(A \Rightarrow B)$ .

In the end, we end up with a list of rules  $A \Rightarrow B$  such that the two conditions above are met.

# Pros and Cons

- Pro: Simplicity and interpretability

# Pros and Cons

- Pro: Simplicity and interpretability
- Pro: Can be run entirely in-database, highly scalable.

# Pros and Cons

- Pro: Simplicity and interpretability
- Pro: Can be run entirely in-database, highly scalable.
- Con: Items with small prevalence are dropped early and not considered.

# Pros and Cons

- Pro: Simplicity and interpretability
- Pro: Can be run entirely in-database, highly scalable.
- Con: Items with small prevalence are dropped early and not considered.
- Con: Will miss strong signals with infrequent data.



# The Setup

This section is taken from Chapter 3 the book *Mining Massive Datasets* by Leskovec, Rajaraman, and Ullman.

# The Setup

This section is taken from Chapter 3 the book *Mining Massive Datasets* by Leskovec, Rajaraman, and Ullman.

The basic idea is to find similar documents in a corpus given an input document.

# The Setup

This section is taken from Chapter 3 the book *Mining Massive Datasets* by Leskovec, Rajaraman, and Ullman.

The basic idea is to find similar documents in a corpus given an input document.

Jaccard similarity: Given finite sets  $A$  and  $B$ ,

$$\text{Sim}_J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

## The Setup

This section is taken from Chapter 3 the book *Mining Massive Datasets* by Leskovec, Rajaraman, and Ullman.

The basic idea is to find similar documents in a corpus given an input document.

Jaccard similarity: Given finite sets  $A$  and  $B$ ,

$$\text{Sim}_J(A, B) = \frac{|A \cap B|}{|A \cup B|}$$

Sometimes we also see this expressed as a "distance":

$$d_J(A, B) = 1 - \text{Sim}_J(A, B)$$

# Shingling

A  $k$ -shingle of a document is any (consecutive) substring of the document.

# Shingling

A  $k$ -shingle of a document is any (consecutive) substring of the document.

Example: If the document in question is "*abcde*", the 2-shingles are  $\{ "ab", "bc", "cd", "de" \}$ .

# Shingling

A  $k$ -shingle of a document is any (consecutive) substring of the document.

Example: If the document in question is "*abcde*", the 2-shingles are  $\{ "ab", "bc", "cd", "de" \}$ .

Another example: If the document in question is "*abracadabra*", the 2-singles are

$$\{ "ab", "br", "ra", "ac", "ca", "da" \}$$

# Shingling

A  $k$ -shingle of a document is any (consecutive) substring of the document.

Example: If the document in question is "*abcde*", the 2-shingles are  $\{ "ab", "bc", "cd", "de" \}$ .

Another example: If the document in question is "*abracadabra*", the 2-singles are

$$\{ "ab", "br", "ra", "ac", "ca", "da" \}$$

There is no sure, fast rule to how long shingles should be, but values between 5 and 9 are common depending on the typical size of the documents (emails vs. websites vs. research papers)



# Hashing

Next we can hash shingles to an integer (often a 32-bit int suffices).

# Hashing

Next we can hash shingles to an integer (often a 32-bit int suffices).

There are alternative methods to shingling utilizing word tokenization and similar techniques, but hashing still plays the pivotal role.

# Characteristic Matrices

Let  $S$  be the universal set of all shingles and let  $S_1, \dots, S_N$  denote the sets of shingles for each document.

# Characteristic Matrices

Let  $S$  be the universal set of all shingles and let  $S_1, \dots, S_N$  denote the sets of shingles for each document.

Let  $M$  denote the matrix such that

$$(m)_{ij} = \begin{cases} 1 & \text{if the } i\text{th shingle is in document } j \\ 0 & \text{otherwise} \end{cases}$$

Note here that the *columns* correspond to the samples in this case.

## Characteristic Matrices

Let  $S$  be the universal set of all shingles and let  $S_1, \dots, S_N$  denote the sets of shingles for each document.

Let  $M$  denote the matrix such that

$$(m)_{ij} = \begin{cases} 1 & \text{if the } i\text{th shingle is in document } j \\ 0 & \text{otherwise} \end{cases}$$

Note here that the *columns* correspond to the samples in this case.

This matrix is called the *characteristic matrix*, and it is typically *very* sparse so it is typical to use a different data structure to represent the matrix.

# Min hashing

- Permute the rows of  $M$ .

# Min hashing

- Permute the rows of  $M$ .
- For each set  $S_j$ , find the first non-zero row.

# Min hashing

- Permute the rows of  $M$ .
- For each set  $S_j$ , find the first non-zero row.
- Do this a few hundred times.



# Min hashing

- Permute the rows of  $M$ .
- For each set  $S_j$ , find the first non-zero row.
- Do this a few hundred times.
- The number of min hashes that coincide for  $S_j$  and  $S_k$  turns out to be  $\text{sim}_J(S_j, S_k)$ .