

Regularization and Logistic Regression

Steve Avsec

Illinois Institute of Technology

January 29, 2024

Overview

- 1 Where We Left Off
- 2 Regularization
- 3 Logistic Regression and the Like

Condition Numbers

$\text{Cond}(A) = \sigma_1(A)/\sigma_r(A)$ where $\sigma_j(A)$ denotes the j th largest singular value of A and $\text{rank}(A) = r$.

Condition Numbers

$\text{Cond}(A) = \sigma_1(A)/\sigma_r(A)$ where $\sigma_j(A)$ denotes the j th largest singular value of A and $\text{rank}(A) = r$.

Alternatively, if $A = U\Sigma V^t$ is the singular value decomposition of A and $A^\dagger = V\Sigma^\dagger U^t$ be the pseudoinverse. Then $\text{Cond}(A) = \sigma_1(A)\sigma_1(A^\dagger)$.

Condition Numbers

$\text{Cond}(A) = \sigma_1(A)/\sigma_r(A)$ where $\sigma_j(A)$ denotes the j th largest singular value of A and $\text{rank}(A) = r$.

Alternatively, if $A = U\Sigma V^t$ is the singular value decomposition of A and $A^\dagger = V\Sigma^\dagger U^t$ be the pseudoinverse. Then $\text{Cond}(A) = \sigma_1(A)\sigma_1(A^\dagger)$.

A matrix A is *ill-conditioned* or *poorly-conditioned* if $\text{Cond}(A)$ is “large” (greater than the precision of the machine).

Least-Squares Regression

Let $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$ be some data ($\mathbf{x}_j \in \mathbb{R}^p$, $y_j \in \mathbb{R}$). Let $f_1, \dots, f_K : \mathbb{R}^p \rightarrow \mathbb{R}$ be a set of “suitable” functions.

Least-Squares Regression

Let $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$ be some data ($\mathbf{x}_j \in \mathbb{R}^p$, $y_j \in \mathbb{R}$). Let $f_1, \dots, f_K : \mathbb{R}^p \rightarrow \mathbb{R}$ be a set of “suitable” functions. Then the linear combination $f = c_1 f_1 + \dots + c_K f_K$ that minimizes

$$\min_{f \in \text{span}(f_1, \dots, f_K)} \|\mathbf{y} - f(X)\|_2^2 = \sum_{j=1}^N |y_j - f(\mathbf{x}_j)|^2$$

can be calculated by $\mathbf{c} = A^\dagger \mathbf{y}$ where

$$A = \begin{bmatrix} f_1(\mathbf{x}_1) & \cdots & f_K(\mathbf{x}_1) \\ \vdots & \vdots & \vdots \\ f_N(\mathbf{x}_N) & \cdots & f_K(\mathbf{x}_N) \end{bmatrix}$$

Least-Squares Regression

Let $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$ be some data ($\mathbf{x}_j \in \mathbb{R}^p$, $y_j \in \mathbb{R}$). Let $f_1, \dots, f_K : \mathbb{R}^p \rightarrow \mathbb{R}$ be a set of “suitable” functions. Then the linear combination $f = c_1 f_1 + \dots + c_K f_K$ that minimizes

$$\min_{f \in \text{span}(f_1, \dots, f_K)} \|\mathbf{y} - f(X)\|_2^2 = \sum_{j=1}^N |y_j - f(\mathbf{x}_j)|^2$$

can be calculated by $\mathbf{c} = A^\dagger \mathbf{y}$ where

$$A = \begin{bmatrix} f_1(\mathbf{x}_1) & \cdots & f_K(\mathbf{x}_1) \\ \vdots & \vdots & \vdots \\ f_1(\mathbf{x}_N) & \cdots & f_K(\mathbf{x}_N) \end{bmatrix}$$

This works great if A is well-conditioned.

Principal Component Regression

What if A is ill-conditioned?

Principal Component Regression

What if A is ill-conditioned?

The most naive solution is to replace A with A_m , the best rank- m approximation to A . This is given by

$$A_m = U\Sigma_m V^t$$

where Σ_m has the same first m values as Σ and sets the other $r - m$ singular values to 0.

Principal Component Regression

What if A is ill-conditioned?

The most naive solution is to replace A with A_m , the best rank- m approximation to A . This is given by

$$A_m = U \Sigma_m V^t$$

where Σ_m has the same first m values as Σ and sets the other $r - m$ singular values to 0.

We can choose m such that $\text{Cond}(A_m)$ is less than a fixed value. For instance, if $\sigma_1(A) = 10^5$ and we wish $\text{Cond}(A_m) < 10^7$, then choose m such that all $\sigma_j(A) < 10^{-2}$ are set to 0.

Ridge Regression

Well we could add a term to our minimizer:

$$\min_{f \in \text{span}(f_1, \dots, f_K)} \|y - \sum_j c_j f_j(X)\|_2 + \lambda \sum |c_j|^2$$

Here λ is a fixed parameter that can be tuned.

Ridge Regression

Well we could add a term to our minimizer:

$$\min_{f \in \text{span}(f_1, \dots, f_K)} \|y - \sum_j c_j f_j(X)\|_2 + \lambda \sum |c_j|^2$$

Here λ is a fixed parameter that can be tuned.
Another way to look at this:

$$\mathbf{c} = (A^t A + \lambda I_K)^{-1} A^t \mathbf{y} = A_\lambda^\dagger$$

where $A_\lambda^\dagger = V \sigma_\lambda^\dagger U^t$ and where σ_λ^\dagger is the following function applied component-wise to Σ .

$$g_\lambda(z) = \frac{z}{z^2 + \lambda}$$

Some Notation

For a vector $\mathbf{x} \in \mathbb{R}^n$, define the p -norm by

$$\|\mathbf{x}\|_p = \left(\sum_{j=1}^n |x_j|^p \right)^{\frac{1}{p}}$$

for $1 \leq p < \infty$. For the ∞ case,

$$\|\mathbf{x}\|_\infty = \max_j |x_j|$$

Similarly for Matrices

For an $N \times K$ matrix A , we can define

$$\|A\|_p = \left(\text{Tr} \left((A^t A)^{\frac{p}{2}} \right) \right)^{\frac{1}{p}}$$

This is equivalent to the vector p -norm of the singular values of A .

Similarly for Matrices

For an $N \times K$ matrix A , we can define

$$\|A\|_p = \left(\text{Tr} \left((A^t A)^{\frac{p}{2}} \right) \right)^{\frac{1}{p}}$$

This is equivalent to the vector p -norm of the singular values of A .

For $p = 2$, this is the *Frobenius* norm of A .

Similarly for Matrices

For an $N \times K$ matrix A , we can define

$$\|A\|_p = \left(\text{Tr} \left((A^t A)^{\frac{p}{2}} \right) \right)^{\frac{1}{p}}$$

This is equivalent to the vector p -norm of the singular values of A .

For $p = 2$, this is the *Frobenius* norm of A .

For $p = \infty$, we recover the "operator norm" of A :

$$\|A\|_\infty = \max_{\mathbf{x} \neq 0} \frac{\|A\mathbf{x}\|_2}{\|\mathbf{x}\|_2}$$

LASSO

Least Absolute Shrinkage and Selection Operator uses the minimizer:

$$\begin{aligned} \min_{f \in \text{span}(f_1, \dots, f_K)} & \|\mathbf{y} - \sum_j c_j f_j(X)\|_2 + \lambda \sum_j |c_j| \\ &= \|\mathbf{y} - \sum_j c_j f_j(X)\|_2 + \lambda \|\mathbf{c}\|_1 \end{aligned}$$

LASSO

Least Absolute Shrinkage and Selection Operator uses the minimizer:

$$\begin{aligned} \min_{f \in \text{span}(f_1, \dots, f_K)} \|\mathbf{y} - \sum_j c_j f_j(X)\|_2 + \lambda \sum_j |c_j| \\ = \|\mathbf{y} - \sum_j c_j f_j(X)\|_2 + \lambda \|\mathbf{c}\|_1 \end{aligned}$$

Because the 1-norm is not differentiable, there is no slick matrix magic we can apply here.

LASSO

Least Absolute Shrinkage and Selection Operator uses the minimizer:

$$\begin{aligned} \min_{f \in \text{span}(f_1, \dots, f_K)} \|\mathbf{y} - \sum_j c_j f_j(X)\|_2 + \lambda \sum_j |c_j| \\ = \|\mathbf{y} - \sum_j c_j f_j(X)\|_2 + \lambda \|\mathbf{c}\|_1 \end{aligned}$$

Because the 1-norm is not differentiable, there is no slick matrix magic we can apply here.

LASSO solutions are sparse (most of the entries of \mathbf{c} are sparse).

The Picture




FIGURE 3.11. Estimation picture for the lasso (left) and ridge regression (right). Shown are contours of the error and constraint functions. The solid blue areas are the constraint regions $|\beta_1| + |\beta_2| \leq t$ and $\beta_1^2 + \beta_2^2 \leq t^2$, respectively, while the red ellipses are the contours of the least squares error function.

Connections to Compressive Sensing

Look for solutions to $A\mathbf{x} = \mathbf{y}$ where A is $k \times N$ for $N \gg k$.

Connections to Compressive Sensing

Look for solutions to $A\mathbf{x} = \mathbf{y}$ where A is $k \times N$ for $N \gg k$.

There are "many" solutions, but "real world" ones tend to be sparse.

Connections to Compressive Sensing

Look for solutions to $A\mathbf{x} = \mathbf{y}$ where A is $k \times N$ for $N \gg k$.

There are "many" solutions, but "real world" ones tend to be sparse.

Reference: *A Mathematical Introduction to Compressive Sensing*, S. Foucart, H. Rauhut

Setup

Recall we could express linear regression by

$$\mathbf{y} = \mathbf{X}\mathbf{c} + \varepsilon$$

Setup

Recall we could express linear regression by

$$\mathbf{y} = X\mathbf{c} + \varepsilon$$

Another way this is commonly expressed is

$$E[\mathbf{y}|X] = X\mathbf{c}$$

Setup

Recall we could express linear regression by

$$\mathbf{y} = X\mathbf{c} + \varepsilon$$

Another way this is commonly expressed is

$$E[\mathbf{y}|X] = X\mathbf{c}$$

This can be modified to:

$$E[\mathbf{y}|X] = g^{-1}(X\mathbf{c})$$

here g is called a link function.

Common Link Functions

- Logit function (Logistic regression)

$$g(z) = \ln \left(\frac{z}{1 - z} \right)$$

- Log function (Poisson regression)

$$g(z) = \ln(z)$$

- Negative inverse (Exponential)

$$g(z) = -\frac{1}{z}$$

- Inverse squared

$$g(z) = \frac{1}{z^2}$$

Gradient Descent: A Detour

Let $F : \mathbb{R}^N \rightarrow \mathbb{R}$ be a (differentiable) function.

Gradient Descent: A Detour

Let $F : \mathbb{R}^N \rightarrow \mathbb{R}$ be a (differentiable) function.

Then at a point $\mathbf{a} \in \mathbb{R}^N$, the function F changes most rapidly in the direction of

$$\pm \nabla F(\mathbf{a})$$

Gradient Descent: A Detour

Let $F : \mathbb{R}^N \rightarrow \mathbb{R}$ be a (differentiable) function.

Then at a point $\mathbf{a} \in \mathbb{R}^N$, the function F changes most rapidly in the direction of

$$\pm \nabla F(\mathbf{a})$$

So if you want to find a local minimum, start at a point \mathbf{a}_0 and then compute the next point by

$$\mathbf{a}_{n+1} = \mathbf{a}_n + \lambda_n \nabla F(\mathbf{a}_n).$$

There are many choices for λ_n , but a common one is

$$\lambda_n = \frac{|(\mathbf{a}_n - \mathbf{a}_{n-1}) \cdot \nabla (F(\mathbf{a}_n) - F(\mathbf{a}_{n-1}))|}{\|\nabla (F(\mathbf{a}_n) - F(\mathbf{a}_{n-1}))\|_2^2}$$

Back to Logistic Regression

Suppose we have $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$ where $\mathbf{x}_j \in \mathbb{R}^k$ and $y_j \in \{0, 1\}$.

Back to Logistic Regression

Suppose we have $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$ where $\mathbf{x}_j \in \mathbb{R}^k$ and $y_j \in \{0, 1\}$.

Objective: Find the $k - 1$ -dimensional hyperplane that best separates the \mathbf{x}_j labeled 0 from those labeled 1.

Back to Logistic Regression

Suppose we have $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$ where $\mathbf{x}_j \in \mathbb{R}^k$ and $y_j \in \{0, 1\}$.

Objective: Find the $k - 1$ -dimensional hyperplane that best separates the \mathbf{x}_j labeled 0 from those labeled 1.

If g is the logit function, then

$$g^{-1}(y) = \frac{1}{1 - e^{-X\mathbf{c}}}$$

Back to Logistic Regression

Suppose we have $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$ where $\mathbf{x}_j \in \mathbb{R}^k$ and $y_j \in \{0, 1\}$.

Objective: Find the $k - 1$ -dimensional hyperplane that best separates the \mathbf{x}_j labeled 0 from those labeled 1.

If g is the logit function, then

$$g^{-1}(y) = \frac{1}{1 - e^{-X\mathbf{c}}}$$

Logistic Regression Finale

So we can express

$$P(y = n|X) = \left(\frac{1}{1 + e^{X\mathbf{c}}} \right)^n \left(1 - \left(\frac{1}{1 + e^{X\mathbf{c}}} \right) \right)^{1-n}$$

Logistic Regression Finale

So we can express

$$P(y = n|X) = \left(\frac{1}{1 - e^{X\mathbf{c}}} \right)^n \left(1 - \left(\frac{1}{1 - e^{X\mathbf{c}}} \right) \right)^{1-n}$$

Taking logs and averaging over samples, we get the function

$$F(\mathbf{c}) = \frac{1}{N} \sum_{j=1}^N y_j \ln(g^{-1}(\mathbf{x}_j\mathbf{c})) + (1 - y_j) \ln(1 - g^{-1}(\mathbf{x}_j\mathbf{c}))$$

Logistic Regression Finale

So we can express

$$P(y = n|X) = \left(\frac{1}{1 - e^{X\mathbf{c}}} \right)^n \left(1 - \left(\frac{1}{1 - e^{X\mathbf{c}}} \right) \right)^{1-n}$$

Taking logs and averaging over samples, we get the function

$$F(\mathbf{c}) = \frac{1}{N} \sum_{j=1}^N y_j \ln(g^{-1}(\mathbf{x}_j\mathbf{c})) + (1 - y_j) \ln(1 - g^{-1}(\mathbf{x}_j\mathbf{c}))$$

From here, compute the gradient of F and use gradient descent!